Actual evaluation trajectories | fixed | 10x10
Each panel uses its own episode target; selected episodes may have different targets.
Path colors progress from early to late; numbers indicate action steps.

Q-Learning / Best Path

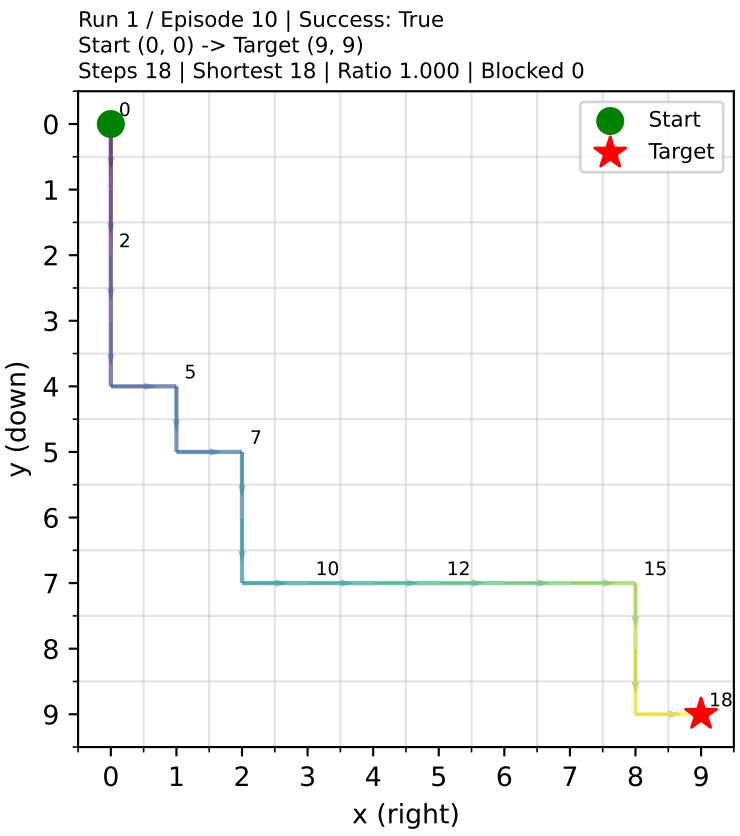

Q-Learning / Representative Path

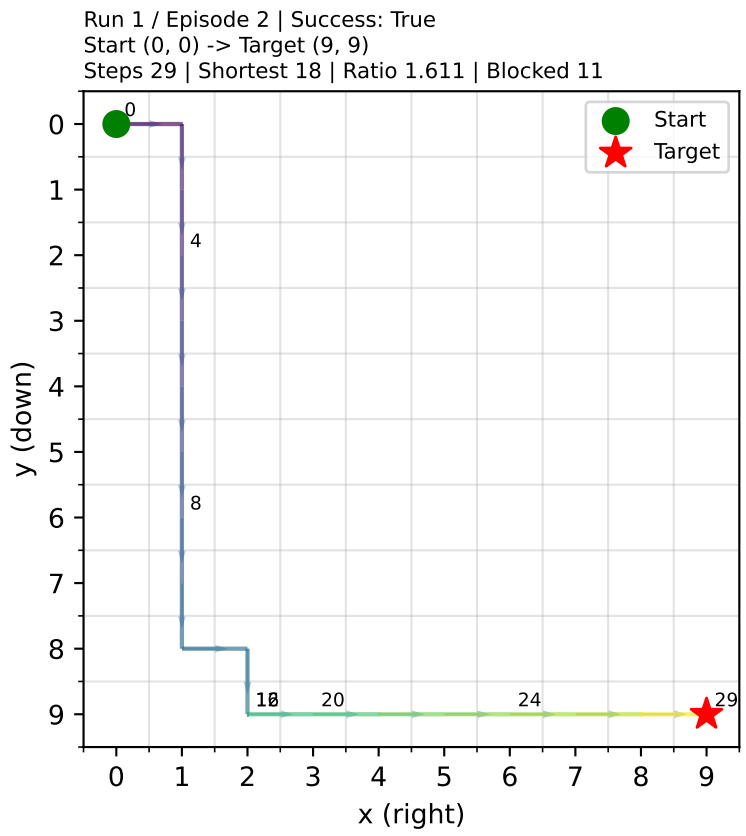

Random Search / Best Path

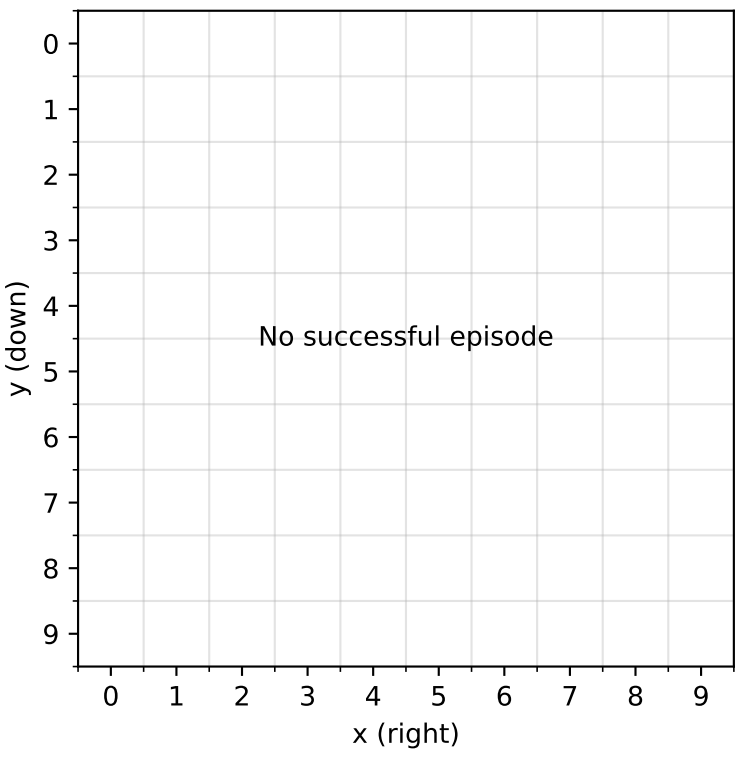

Random Search / Representative Path

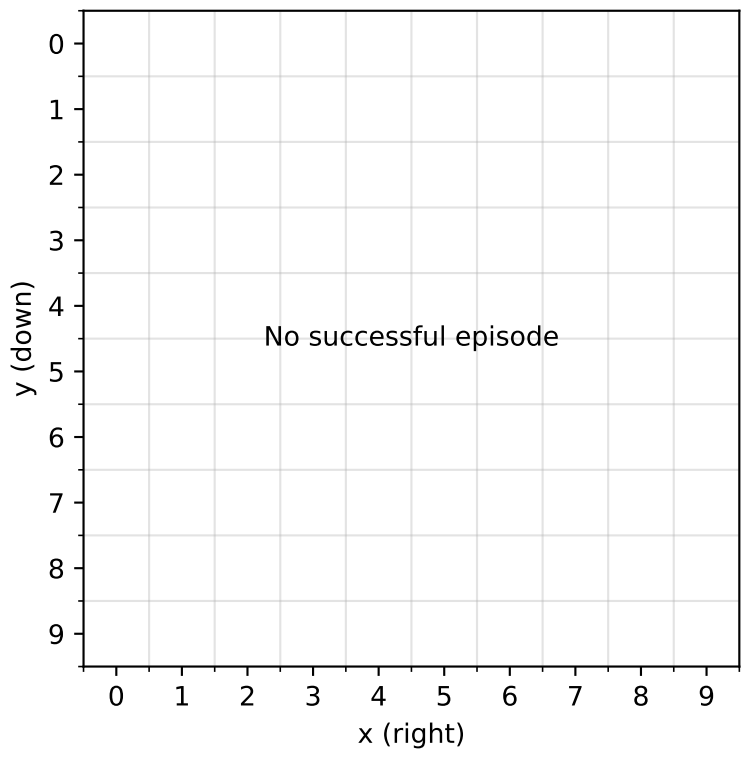

All evaluation episodes pooled across runs (steps and distance include failures):
Q-Learning: success 80.0% | mean steps (all) 44.40 | mean path ratio (success only, positive shortest distance) 1.694 | mean final distance (all, cells) 0.60
Random Search: success 0.0% | mean steps (all) 100.00 | mean path ratio (success only, positive shortest distance) N/A | mean final distance (all, cells) 8.39